



# **Final report**

**Member:**

**Date : 12/18**

# **AD Data**

# Model selection – ResNet34




Solve gradient vanishing/exploding to improve model performance

Figure 1-20. Residual structures


# Training – selection best model

```
# Paths  
str_Workspace = os.getcwd()  
  
# Path related  
str_DatasetPath = str_ADPATH  
  
# Hyperparameters  
int_Epochs = 100  
int_BatchSize = 15  
  
int_ImageSize = 96  
win_size = (int_ImageSize, int_ImageSize, int_ImageSize)  
  
float_LearningRate= 0.01  
float_TestRatio = 0.2  
  
# int_RandSeedModel = 42  
int_RandSeedData = 5
```




# Evaluation metrics – test data

	precision	recall	f1-score	support
Normal	1.00	0.89	0.94	9
AD	0.91	1.00	0.95	10
accuracy			0.95	19
macro avg	0.95	0.94	0.95	19
weighted avg	0.95	0.95	0.95	19




# Interpretability - Patients

```
cam = monai.visualize.GradCAM(nn_module=model, target_layers="layer4.1.conv2")
```




# Interpretability - Normal person


actual: Normal  
pred: Normal  
AD: -13.1  
Normal: 13.0




CAM




Occ. sens.




actual: Normal  
pred: Normal  
AD: -10.8  
Normal: 10.7




CAM




Occ. sens.




actual: Normal  
pred: Normal  
AD: 0.00441  
Normal: 0.0225




CAM




Occ. sens.




actual: Normal  
pred: Normal  
AD: -12.8  
Normal: 12.7




CAM




Occ. sens.




actual: Normal  
pred: Normal  
AD: -9.94  
Normal: 9.94



CAM




Occ. sens.



# **ASD Data**

# Model selection – DenseNet121



- DenseNet121 and ResNet34 demonstrate poor accuracy
- DenseNet121 uses fewer computational resources

# Training – selection best model

```
# Paths
str_Workspace = os.getcwd()

# Path related
str_DatasetPath = str_ADPATH


# Hyperparameters
int_Epochs = 100
int_BatchSize = 15

int_ImageSize = 96
win_size = (int_ImageSize, int_ImageSize, int_ImageSize)

float_LearningRate= 0.01
float_TestRatio = 0.2


# int_RandSeedModel = 42
int_RandSeedData = 5
```

Validation data 56.25% accuracy in epoch 17




# Evaluation metrics – test data

	precision	recall	f1-score	support
Normal	0.57	0.86	0.69	43
ASD	0.60	0.24	0.35	37
accuracy			0.57	80
macro avg	0.58	0.55	0.52	80
weighted avg	0.58	0.57	0.53	80




# Interpretability - Patients


actual: ASD  
pred: ASD  
ASD: 0.099  
Normal: -0.104




CAM




Occ. sens.




actual: ASD  
pred: ASD  
ASD: 0.0481  
Normal: -0.0616




CAM




Occ. sens.




actual: ASD  
pred: ASD  
ASD: 0.0689  
Normal: -0.0822




CAM




Occ. sens.




actual: ASD  
pred: ASD  
ASD: 0.125  
Normal: -0.132




CAM




Occ. sens.





actual: ASD  
pred: ASD  
ASD: 0.0689  
Normal: -0.0822



CAM




Occ. sens.





# Interpretability - Normal person


actual: Normal  
pred: Normal  
ASD: -0.0762  
Normal: 0.0738



CAM





Occ. sens.




actual: Normal  
pred: Normal  
ASD: -0.0705  
Normal: 0.0672



CAM




Occ. sens.





actual: Normal  
pred: Normal  
ASD: -0.0666  
Normal: 0.0612




CAM




Occ. sens.




actual: Normal  
pred: Normal  
ASD: -0.0442  
Normal: 0.0343




CAM





Occ. sens.





actual: Normal  
pred: Normal  
ASD: -0.0666  
Normal: 0.0612



CAM




Occ. sens.



**Bonus –**  
**three-class classification**  
**or multimodel**


# Model selection – Variational autoencoder




$$Loss = \sum x \log x' + (1 - x) \log(1 - x') - \frac{1}{2} \sum (1 + \log(\sigma_i^2) - m_i^2 - \sigma_i^2)$$

# Model selection – Variational autoencoder

AE:



VAE:



Learning progressive disease  
to extract latent representation

# Model architecture

```
# Create DenseNet121, CrossEntropyLoss and Adam optimizer
NetArch = monai.networks.nets.VarAutoEncoder(
    spatial_dims=3,
    in_shape = (1, 96, 96, 96),
    out_channels= 1,
    latent_size = 2, Latent representation size
    channels=(4, 4, 4),
    strides=(1, 2, 1),
    num_res_units = 1
)
model=NetArch.to(device)
```

# Model architecture

## Encoder:


Layer (type)	Output Shape	Param #
Conv3d-1	[ -1, 4, 96, 96, 96]	8
Conv3d-2	[ -1, 4, 96, 96, 96]	112
InstanceNorm3d-3	[ -1, 4, 96, 96, 96]	0
PReLU-4	[ -1, 4, 96, 96, 96]	1
ResidualUnit-5	[ -1, 4, 96, 96, 96]	0
Conv3d-6	[ -1, 4, 48, 48, 48]	436
Conv3d-7	[ -1, 4, 48, 48, 48]	436
InstanceNorm3d-8	[ -1, 4, 48, 48, 48]	0
PReLU-9	[ -1, 4, 48, 48, 48]	1
ResidualUnit-10	[ -1, 4, 48, 48, 48]	0
Identity-11	[ -1, 4, 48, 48, 48]	0
Conv3d-12	[ -1, 4, 48, 48, 48]	436
InstanceNorm3d-13	[ -1, 4, 48, 48, 48]	0
PReLU-14	[ -1, 4, 48, 48, 48]	1
ResidualUnit-15	[ -1, 4, 48, 48, 48]	0
Identity-16	[ -1, 4, 48, 48, 48]	0
Linear-17	[ -1, 2]	884,738
Linear-18	[ -1, 2]	884,738
Linear-19	[ -1, 442368]	1,327,104

Latent representation

## Decoder:

ConvTranspose3d-20	[ -1, 4, 48, 48, 48]	436
InstanceNorm3d-21	[ -1, 4, 48, 48, 48]	0
PReLU-22	[ -1, 4, 48, 48, 48]	1
Identity-23	[ -1, 4, 48, 48, 48]	0
Conv3d-24	[ -1, 4, 48, 48, 48]	436
InstanceNorm3d-25	[ -1, 4, 48, 48, 48]	0
PReLU-26	[ -1, 4, 48, 48, 48]	1
ResidualUnit-27	[ -1, 4, 48, 48, 48]	0
ConvTranspose3d-28	[ -1, 4, 96, 96, 96]	436
InstanceNorm3d-29	[ -1, 4, 96, 96, 96]	0
PReLU-30	[ -1, 4, 96, 96, 96]	1
Identity-31	[ -1, 4, 96, 96, 96]	0
Conv3d-32	[ -1, 4, 96, 96, 96]	436
InstanceNorm3d-33	[ -1, 4, 96, 96, 96]	0
PReLU-34	[ -1, 4, 96, 96, 96]	1
ResidualUnit-35	[ -1, 4, 96, 96, 96]	0
ConvTranspose3d-36	[ -1, 1, 96, 96, 96]	109
InstanceNorm3d-37	[ -1, 1, 96, 96, 96]	0
PReLU-38	[ -1, 1, 96, 96, 96]	1
Identity-39	[ -1, 1, 96, 96, 96]	0
Conv3d-40	[ -1, 1, 96, 96, 96]	28
ResidualUnit-41	[ -1, 1, 96, 96, 96]	0

# Pipeline



# Training model

```
# Hyperparameters
int_Epochs = 10
int_BatchSize = 8

int_ImageSize = 96
win_size = (int_ImageSize, int_ImageSize, int_ImageSize)

float_LearningRate= 1e-6
float_TestRatio = 0.2
```

```
def loss_function(recon_x, x, mu, log_var, beta):
    mse_loss = F.mse_loss(recon_x, x)
    kld = -0.5 * beta * torch.sum(1 + log_var - mu.pow(2) - log_var.exp())
    return mse_loss + kld

optimizer = torch.optim.Adam(model.parameters(), float_LearningRate)
```

```
mu_list = []
labels_list = []

for test_data in test_loader:
    test_images = test_data[0].to(device)
    labels = test_data[1].to(device)



    _, mu, _, _ = model(test_images)

    mu_list.append(mu.cpu().detach().numpy())
    labels_list.append(labels.cpu().detach().numpy())


mu_array = np.concatenate(mu_list, axis=0)

labels_array = np.concatenate(labels_list, axis=0)
labels_array = np.argmax(labels_array, axis=1)
```

# Analyze latent representations



# Embeddings(mu) - PCA



# Machine learning performance (LDA)

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
X = mu_df

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda = LinearDiscriminantAnalysis(n_components=2)

from sklearn.model_selection import cross_val_score

cv_scores = cross_val_score(lda, X, y, cv=4, scoring='accuracy')
print(f"Cross-validation scores: {cv_scores}")
print(f"Average accuracy: {cv_scores.mean():.4f}")
```

```
Cross-validation scores: [0.42857143 0.85714286 0.57142857 0.42857143]
Average accuracy: 0.5714
```

# Multimodal (LDA)

```
df_test_drop = df_test[['Sex', 'Age', 'ADAS', 'MMSE', 'NPI']].copy()
df_test_drop['Sex'] = df_test_drop['Sex'].map({'M': 0, 'F': 1})
df_test_drop = df_test_drop.fillna(df_test_drop.median())
Cross-validation scores: [0.85714286 0.71428571 0.71428571 0.85714286]
Average accuracy: 0.7857
```

## Concat latent representation

```
mu_df = pd.DataFrame(mu_array, columns=['Feature1', 'Feature2'])
df_combined = pd.concat([df_test_drop, mu_df], axis=1)
```

```
Cross-validation scores: [1.          0.71428571 0.85714286 0.71428571]
Average accuracy: 0.8214
```